

## RAG PROTOTYPE- PLAN

### 1. Retrieval Index Construction

#### 1.1 Vector Database Implementation

1. Technology: FAISS (Facebook AI Similarity Search)
2. Embedding Model: nomic-embed-text served through Ollama
3. Persistant Storage: faiss\_index/ directory
4. Index Type: Flat L2 index for exact similarity search

#### 1.2 Construction Process

1. Load PDF: Extract text from pdf using PyMuPDFLoader
2. Split Documents: Break into chunks using RecursiveCharacterTextSplitter
3. Generate Embeddings: Convert each chunk to vectors using embedding model
4. Build Index: Create FAISS index from document embeddings
5. Persist: Save index to disk at faiss\_index/ for reuse

#### 1.3 Retrieval Mechanism

- Returns top-k most relevant document chunks for each query
- Preserves metadata including page numbers
- Loads persisted index on startup
- Converts queries to embeddings for similarity search

### 2. Chunking Strategy

#### 2.1 Configuration Parameters

- Chunk Size: 300 characters
- Chunk Overlap: 25 characters
- Splitter Type: RecursiveCharacterTextSplitter
- Separator Hierarchy: Paragraph, Line, Space, Character
- Source document reference, Page number, Chunk content with formatting is shown

### 3. Prompt Template Family for Audience Adaptation

**Standard RAG Query:** Use retrieved context, cite [Source X], format equations in LaTeX, avoid HTML tags, and admit when the answer is not in the context.

**Question Contextualization:** Reformulate follow-up into standalone by referencing chat history when needed, leave them unchanged if self-contained, and don't attempt to answer

**Refinement Query:** Review the previous output alongside feedback, enhance response additional context, return JSON with changes and reasons